

Module 3

Thursday, July 03, 2025 1:17 AM

What is classification?




Supervised ML method

Uses fully trained models to predict labels on new data

Labels form a categorical variable with discrete values

What is supervised learning?



Understands data in context when answering a question

Ensures accuracy in predictions

Model adjusts the data to fit the algorithm and classifies it accordingly

Applications of classification

| | tenure | age | address | income | ed | employ | equip | callcard | wireless | churn |
|---|--------|------|---------|--------|-----|--------|-------|----------|----------|-------|
| 0 | 11.0 | 33.0 | 7.0 | 136.0 | 5.0 | 5.0 | 0.0 | 1.0 | 1.0 | Yes |
| 1 | 33.0 | 33.0 | 12.0 | 33.0 | 2.0 | 0.0 | 0.0 | 0.0 | 0.0 | Yes |
| 2 | 23.0 | 30.0 | 9.0 | 30.0 | 1.0 | 2.0 | 0.0 | 0.0 | 0.0 | No |
| 3 | 38.0 | 35.0 | 5.0 | 76.0 | 2.0 | 10.0 | 1.0 | 1.0 | 1.0 | No |
| 4 | 7.0 | 35.0 | 14.0 | 80.0 | 2.0 | 15.0 | 0.0 | 1.0 | 0.0 | ? |

Churn prediction: If customer will discontinue service

Customer segmentation: Predict the category of a customer

Advertising: Predict if a customer will respond to a campaign

Use cases of classification

Loan default prediction

| age | ed | employ | address | income | debtinc | creddebt | othdebt | default |
|-----|----|--------|---------|--------|---------|----------|---------|---------|
| 41 | 3 | 17 | 12 | 176 | 9.3 | 11.359 | 5.009 | 1 |
| 27 | 1 | 10 | 6 | 31 | 17.3 | 1.362 | 4.001 | 0 |
| 40 | 1 | 15 | 14 | 55 | 5.5 | 0.856 | 2.169 | 0 |
| 41 | 1 | 15 | 14 | 120 | 2.9 | 2.659 | 0.821 | 0 |
| 24 | 2 | 2 | 0 | 28 | 17.3 | 1.787 | 3.057 | 1 |
| 41 | 2 | 5 | 5 | 25 | 10.2 | 0.393 | 2.157 | 0 |
| 39 | 1 | 20 | 9 | 67 | 30.6 | 3.834 | 16.668 | 0 |
| 43 | 1 | 12 | 11 | 38 | 3.6 | 0.129 | 1.239 | 0 |
| 24 | 1 | 3 | 4 | 19 | 24.4 | 1.358 | 3.278 | 1 |
| 36 | 1 | 0 | 13 | 25 | 19.7 | 2.778 | 2.147 | 0 |



| age | ed | employ | address | income | debtinc | creddebt | othdebt | default |
|-----|----|--------|---------|--------|---------|----------|---------|---------|
| 37 | 2 | 16 | 10 | 130 | 9.3 | 10.23 | 3.21 | |

Prediction

Trained Classifier

Use cases of classification

Multiclass drug prescription

| Age | Sex | BP | Cholesterol | Na | K | Drug |
|-----|-----|--------|-------------|-------|-------|-------|
| 23 | F | HIGH | HIGH | 0.793 | 0.031 | drugY |
| 47 | M | LOW | HIGH | 0.739 | 0.056 | drugC |
| 47 | M | LOW | HIGH | 0.697 | 0.069 | drugC |
| 28 | F | NORMAL | HIGH | 0.564 | 0.072 | drugX |
| 61 | F | LOW | HIGH | 0.559 | 0.031 | drugY |
| 22 | F | NORMAL | HIGH | 0.677 | 0.079 | drugX |
| 49 | F | NORMAL | HIGH | 0.79 | 0.049 | drugY |
| 41 | M | LOW | HIGH | 0.767 | 0.069 | drugC |
| 60 | M | NORMAL | HIGH | 0.777 | 0.051 | drugY |
| 43 | M | LOW | NORMAL | 0.526 | 0.027 | drugY |

Categorical Variable

Modeling

Prediction

Classifier

Predicted Labels


Classification algorithms



Naïve Bayes



Logistic regression



Decision trees




One-versus-all strategy


- Binary classifier:** One for each class label
- Assigned a single label that defines target class
- Task:** Binary prediction for every data point for a one-versus-the-rest classifier
- K-classes** = K binary classifiers




One-versus-all strategy



One-versus-one strategy




One-versus-one strategy



Decision tree

- Each internal node corresponds to a test
- Each branch corresponds to the result of the test
- Each terminal, or leaf node, assigns its data to a class



How to build a decision tree?

| Patient ID | Age | Sex | BP | Cholesterol | Drug |
|------------|------------|-----|--------|-------------|--------|
| p1 | Young | F | High | Normal | Drug A |
| p2 | Young | F | High | High | Drug A |
| p3 | Middle-age | F | High | Normal | Drug B |
| p4 | Senior | F | Normal | Normal | Drug B |
| p5 | Senior | M | Low | Normal | Drug B |
| p6 | Senior | M | Low | High | Drug A |
| p7 | Middle-age | M | Low | High | Drug B |
| p8 | Young | F | Normal | Normal | Drug A |
| p9 | Young | M | Low | Normal | Drug B |
| p10 | Senior | M | Normal | Normal | Drug B |
| p11 | Young | M | Normal | High | Drug B |
| p12 | Middle-age | F | Normal | High | Drug B |




- Consider features of a data set
- Example in medical study: Age, sex, blood pressure, and cholesterol
- Use training part of data set to build decision tree
- Use decision tree to predict class of unknown

| | | | | | |
|-----|------------|---|--------|--------|--------|
| p13 | Middle-age | M | High | Normal | Drug B |
| p14 | Senior | F | Normal | High | Drug A |
| p15 | Middle-age | F | Low | Normal | ? |


patient

Prediction

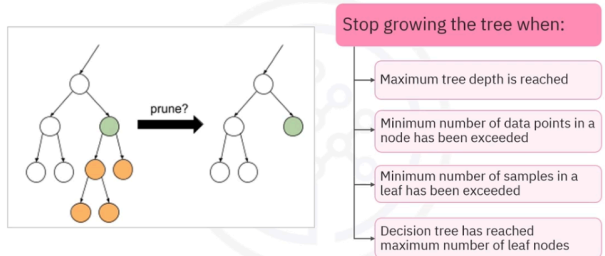
Patient classifier example




Patient classifier example



Tree pruning



Pruning decision tree example



Module 3 Summary and Highlights

Congratulations! You have completed this lesson. At this point in the course, you know:

- Classification is a supervised machine learning method used to predict labels on new data with applications in churn prediction, customer segmentation, loan default prediction, and multiclass drug prescriptions.
- Binary classifiers can be extended to multiclass classification using one-versus-all or one-versus-one strategies.
- A decision tree classifies data by testing features at each node, branching based on test results, and assigning classes at leaf nodes.
- Decision tree training involves selecting features that best split the data and pruning the tree to avoid overfitting.
- Information gain and Gini impurity are used to measure the quality of splits in decision trees.
- Regression trees are similar to decision trees but predict continuous values by recursively splitting data to maximize information gain.
- Mean Squared Error (MSE) is used to measure split quality in regression trees.
- K-Nearest Neighbors (k-NN) is a supervised algorithm used for classification and regression by assigning labels based on the closest labeled data points.

- To optimize k-NN, test various k values and measure accuracy, considering class distribution and feature relevance.
- Support Vector Machines (SVM) build classifiers by finding a hyperplane that maximizes the margin between two classes, effective in high-dimensional spaces but sensitive to noise and large datasets.
- The bias-variance tradeoff affects model accuracy, and methods such as bagging, boosting, and random forests help manage bias and variance to improve model performance.
- Random forests use bagging to train multiple decision trees on bootstrapped data, improving accuracy by reducing variance.

Cheat Sheet: Building Supervised Learning Models

Common supervised learning models

| Process Name | Brief Description | Code Syntax |
|---|--|--|
| One vs One classifier (using logistic regression) | <p>Process: This method trains one classifier for each pair of classes.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'estimator': Base classifier (e.g., logistic regression) <p>Pros: Can work well for small datasets.</p> <p>Cons: Computationally expensive for large datasets.</p> <p>Common applications: Multiclass classification problems where the number of classes is relatively small.</p> | <pre>from sklearn.multiclass import OneVsOneClassifier from sklearn.linear_model import LogisticRegression model = OneVsOneClassifier(Logisti cRegression())</pre> |
| One vs All classifier (using logistic regression) | <p>Process: Trains one classifier per class, where each classifier distinguishes between one class and the rest.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'estimator': Base classifier (e.g., Logistic Regression) - 'multi_class': Strategy to handle multiclass classification ('ovr') <p>Pros: Simpler and more scalable than One vs One.</p> <p>Cons: Less accurate for highly imbalanced classes.</p> <p>Common applications: Common in multiclass classification problems such as image classification.</p> | <pre>from sklearn.multiclass import OneVsRestClassifier from sklearn.linear_model import LogisticRegression model = OneVsRestClassifier(Logist icRegression()) or from sklearn.linear_model import LogisticRegression model_ova = LogisticRegression(mu lti_class='ovr')</pre> |
| Decision tree classifier | <p>Process: A tree-based classifier that splits data into smaller subsets based on feature values.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'max_depth': Maximum depth of the tree <p>Pros: Easy to interpret and</p> | <pre>from sklearn.tree import DecisionTreeClassifier model = DecisionTreeClassifier(max _depth=5)</pre> |

| | | |
|--------------------------------|---|---|
| | <p>visualize.</p> <p>Cons: Prone to overfitting if not pruned properly.</p> <p>Common applications: Classification tasks, such as credit risk assessment.</p> | |
| Decision tree regressor | <p>Process: Similar to the decision tree classifier, but used for regression tasks to predict continuous values.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'max_depth': Maximum depth of the tree <p>Pros: Easy to interpret, handles nonlinear data.</p> <p>Cons: Can overfit and perform poorly on noisy data.</p> <p>Common applications: Regression tasks, such as predicting housing prices.</p> | <pre>from sklearn.tree import DecisionTreeRegressor model = DecisionTreeRegressor(max_ depth=5)</pre> |
| Linear SVM classifier | <p>Process: A linear classifier that finds the optimal hyperplane separating classes with a maximum margin.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'C': Regularization parameter - 'kernel': Type of kernel function ('linear', 'poly', 'rbf', etc.) - 'gamma': Kernel coefficient (only for 'rbf', 'poly', etc.) <p>Pros: Effective for high-dimensional spaces.</p> <p>Cons: Not ideal for nonlinear problems without kernel tricks.</p> <p>Common applications: Text classification and image recognition.</p> | <pre>from sklearn.svm import SVC model = SVC(kernel='linear', C=1.0)</pre> |
| K-nearest neighbors classifier | <p>Process: Classifies data based on the majority class of its nearest neighbors.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'n_neighbors': Number of neighbors to use - 'weights': Weight function used in prediction ('uniform' or 'distance') - 'algorithm': Algorithm used to compute the nearest neighbors ('auto', 'ball_tree', 'kd_tree', 'brute') <p>Pros: Simple and effective for small datasets.</p> <p>Cons: Computationally expensive as the dataset grows.</p> <p>Common applications: Recommendation systems, image recognition.</p> | <pre>from sklearn.neighbors import KNeighborsClassifier model = KNeighborsClassifier(n_nei ghbors=5, weights='uniform')</pre> |
| Random Forest regressor | <p>Process: An ensemble method using multiple decision trees to improve accuracy and reduce overfitting.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'n_estimators': Number of trees in the forest - 'max_depth': Maximum depth of each tree <p>Pros: Less prone to overfitting than individual decision trees.</p> <p>Cons: Model complexity increases with the number of trees.</p> <p>Common applications: Regression tasks such as predicting sales or stock prices.</p> | <pre>from sklearn.ensemble import RandomForestRegressor model = RandomForestRegressor(n_es timators=100, max_depth=5)</pre> |

| | | |
|-------------------|--|---|
| XGBoost regressor | <p>Process: A gradient boosting method that builds trees sequentially to correct errors from previous trees.</p> <p>Key hyperparameters:</p> <ul style="list-style-type: none"> - 'n_estimators': Number of boosting rounds - 'learning_rate': Step size to improve accuracy - 'max_depth': Maximum depth of each tree <p>Pros: High accuracy and works well with large datasets.</p> <p>Cons: Computationally intensive, complex to tune.</p> <p>Common applications: Predictive modeling, especially in Kaggle competitions.</p> | <pre>import xgboost as xgb model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1, max_depth=5)</pre> |
|-------------------|--|---|

Associated functions used

| Method Name | Brief Description | Code Syntax |
|----------------|---|--|
| OneHotEncoder | Transforms categorical features into a one-hot encoded matrix. | <pre>from sklearn.preprocessing import OneHotEncoder encoder = OneHotEncoder(sparse=False) encoded_data = encoder.fit_transform(c ategorical_data)</pre> |
| accuracy_score | Computes the accuracy of a classifier by comparing predicted and true labels. | <pre>from sklearn.metrics import accuracy_score accuracy = accuracy_score(y_true, y_pred)</pre> |
| LabelEncoder | Encodes labels (target variable) into numeric format. | <pre>from sklearn.preprocessing import LabelEncoder encoder = LabelEncoder() encoded_labels = encoder.fit_transform(l abels)</pre> |

| | | |
|-----------------------|---|--|
| | | |
| plot_tree | Plots a decision tree model for visualization. | <pre>from sklearn.tree import plot_tree plot_tree(model, max_depth=3, filled=True)</pre> |
| normalize | Scales each feature to have zero mean and unit variance (standardization). | <pre>from sklearn.preprocessing import normalize normalized_data = normalize(data, norm='l2')</pre> |
| compute_sample_weight | Computes sample weights for imbalanced datasets. | <pre>from sklearn.utils.class_wi ght import compute_sample_weight weights = compute_sample_weight(c lass_weight='balanced', y=y)</pre> |
| roc_auc_score | Computes the Area Under the Receiver Operating Characteristic Curve (AUC-ROC) for binary classification models. | <pre>from sklearn.metrics import roc_auc_score auc = roc_auc_score(y_true, y_score)</pre> |

